

PRODUCT STRATEGY REVIEW

Job Tracker

From a personal automation to a product strangers can use — the problem, the original build, the plan, and what we actually settled on.

The Problem Statement

1

Too many applications, too little memory

Applying to 30–50 jobs a day means every posting needs a freshly tailored CV and cover letter — and it's easy to lose track of what was sent where.

2

Tracking becomes its own chore

Status, dates and follow-ups end up scattered across an inbox, a spreadsheet, and memory — with no single place that shows the whole pipeline.

3

Constant app-switching adds cognitive load

Tailoring a CV, writing a cover letter, and logging the application each live in a different tool, so every single application means several context switches.

4

The pivot: from personal tool to shared product

Turn this from a one-person script into something real strangers can use — each bringing their own free AI key, with no security compromise and no cost to the builder.

Where We Started: The Original Build

A personal tool, built for one user

- Gmail scanning via GitHub Actions — reads confirmation emails to log applications automatically
- Google Sheet as the database — zero infrastructure, built around a single account
- FastAPI webapp on Render's free tier — tailors CVs and cover letters using the owner's own Gemini key
- One shared password (HTTP Basic Auth) — fine for one person, not for the public

30–50

applications per day

The personal usage volume that shaped the original design — and the reason tailoring and tracking needed to be fast, not just functional.

What We Planned To Do

THE ASK

- Open the tool to real strangers, not just the builder
- Each user brings their own free Gemini API key
- The result must be “completely secure”
- Zero cost to the builder, under any circumstances

AN IDEA WE EXPLORED

REJECTED

Watch the live application form

An embedded browser that auto-captures job title and company while the user fills out an application — so nothing has to be typed twice.

- Can't selectively see 3 fields — it's all-or-nothing page access
- Most major ATS platforms forbid automated form interaction
- Needs paid, per-user infrastructure — breaks “zero cost”
- Looks like a keylogger from the outside, however it's built

What We Settled On



Stateless tailoring

Paste a job link or text, get a tailored CV + cover letter — nothing is stored on the server.



Paste, don't automate

No Gmail scanning, no embedded browser for v1 — the user pastes the job info by hand.



Accounts via Firebase

Firebase Auth for login, Firestore for storage — each user's data isolated by security rules.



Pipeline board

A HubSpot-style board — drag each application through stages like Applied, Interview, Offer.



Downloadable spreadsheet

Export the full tracker as a file on demand, generated fresh each time — nothing persisted in between.



Telegram updates

Users connect their own bot token and chat ID to get notified of status changes.

The Trade-offs

Decision	What We Gained	What We Gave Up
Stateless tailoring tool	Nothing to breach — zero server-side liability	No memory between applications or tailoring history
No embedded browser	Avoids ToS violations and the “keylogger” trust problem	Users still paste job details in by hand
Bring-your-own Gemini key	Zero AI cost to the builder, at any volume	Key handling becomes the user's own responsibility
Accounts via Firebase	Enables the pipeline board, Telegram, and history	Reopens storage and security questions; no longer anonymous
Free-tier infrastructure	The entire product costs \$0 to run	Real limits under concurrent strangers — still unverified at scale

Where This Leaves Us

Scope for v1 is locked: stateless tailoring, paste-only input, Firebase accounts, the pipeline board, spreadsheet export, and Telegram updates.

THE ONE THING TO DO FIRST

Verify Gemini's free-tier rate limits under simulated concurrent, multi-tenant load — the entire zero-cost pitch depends on this holding up.